

DATA 607 Project 2

Samuel C

2025-03-09

Overview

In this project I have chosen three datasets provided from my classmates to tidy and clean. These datasets involve sales data, weather data across different cities, and lastly emissions data across different countries. Not only will these datasets be tidied, I will also perform a bit of exploratory data analysis in order to consider possible relationships among the data in each dataset.

Getting Started

First, we must load the packages and data we will use. I have stored the data on my github across three separate .csv files.

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr    1.5.1
## v ggplot2     3.5.1      v tibble     3.2.1
## v lubridate  1.9.4      v tidyr      1.3.1
## v purrr       1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
untidy_emissions <- read.csv("https://raw.githubusercontent.com/scrummett/DATA607/refs/heads/main/TotalEmissions.csv")
untidy_sales <- read.csv("https://raw.githubusercontent.com/scrummett/DATA607/refs/heads/main/salesdata.csv")
untidy_weather <- read.csv("https://raw.githubusercontent.com/scrummett/DATA607/refs/heads/main/weatherdata.csv")
```

Now with the data loaded, I will begin tidying and EDA from the easiest to the most intensive.

Sales Data

```
head(untidy_sales)
```

```
##   Product.Name Region Jan.Sales Feb.Sales Mar.Sales Apr.Sales May.Sales
## 1 Product A   North      100      110      120      130      140
## 2 Product A   South      200      210      220      230      240
## 3 Product A   East       300      310      320      330      340
## 4 Product B   North      150      160      170      180      190
## 5 Product B   South      250      260      270      280      290
```

```
## 6      Product B      East      350      360      370      380      390
## Jun.Sales
## 1        150
## 2        250
## 3        350
## 4        200
## 5        300
## 6        400
```

This dataset has sales per month separated out across different columns, however a clean version of this dataset would have “months” be a column itself, and total sales figures being a separate column as well.

```
untidy_sales <- untidy_sales |>
  pivot_longer(
    cols = ends_with(".Sales"),
    names_to = "Month",
    values_to = "Sales")
head(untidy_sales)
```

```
## # A tibble: 6 x 4
##   Product.Name Region Month      Sales
##   <chr>         <chr> <chr>    <int>
## 1 Product A     North Jan.Sales  100
## 2 Product A     North Feb.Sales  110
## 3 Product A     North Mar.Sales  120
## 4 Product A     North Apr.Sales  130
## 5 Product A     North May.Sales  140
## 6 Product A     North Jun.Sales  150
```


We now have a “tidy” dataset, however the data in the columns can be cleaned up to avoid redundancy.

```
untidy_sales <- untidy_sales |>
  mutate(Month = str_remove(Month, ".Sales"),
         Product.Name = str_remove(Product.Name, "Product "))
tidy_sales <- untidy_sales |>
  rename("Product" = "Product.Name")
head(tidy_sales)
```


```
## # A tibble: 6 x 4
##   Product Region Month Sales
##   <chr>    <chr> <chr> <int>
## 1 A      North Jan    100
## 2 A      North Feb    110
## 3 A      North Mar    120
## 4 A      North Apr    130
## 5 A      North May    140
## 6 A      North Jun    150
```

After cleaning up observations and changing the title of a column, we now have a tidy dataset of sales data.

```
tidy_sales |>  
  ggplot(aes(x = Product, y = Sales)) +  
  geom_boxplot()
```




```
tidy_sales |>  
  ggplot(aes(x = Region, y = Sales)) +  
  geom_boxplot()
```



Here we have graphs showing total sales by product and then by region. The graphs show that while products A and B sell similar albeit different amounts, sales of product C lag behind severely. Additionally, the North has far fewer sales than both the South and the East.

One question posed in our discussion forum asked about sales performance by product across regions, which can be seen in the following graph.

```
tidy_sales |>  
  ggplot(aes(x = Region, y = Sales, fill = Product)) +  
  geom_boxplot()
```



From this we can see that product B sells the most across every region, and that product C sells the least.

Weather

While our sales data only had months and sales to elongate, the dataset for weather has one more.

```
head(untidy_weather)
```

```
##           City Temp_Jan Temp_Feb Temp_Mar Humid_Jan Humid_Feb Humid_Mar
## 1   New York   32°F     35°F     42°F     75%      72%      68%
## 2 Los Angeles  58°F     60°F     65°F     65%      63%      60%
## 3   Chicago   28°F     30°F     40°F     80%      78%      75%
```

Again, months are spread across the data as separate variables, however we also have temperature and humidity to separate out into individual columns as well.

```
untidy_weather_left <- untidy_weather |>
  pivot_longer(
    cols = starts_with("Temp_"),
    names_to = c("Month"),
    names_prefix = "Temp_",
    values_to = "Temp_F")
untidy_weather_right <- untidy_weather |>
  pivot_longer(
    cols = starts_with("Humid_"),
    names_to = "Month",
```

```

names_prefix = "Humid_",
values_to = "Humidity_Percent"
)
head(untidy_weather_left)

```

```

## # A tibble: 6 x 6
##   City      Humid_Jan Humid_Feb Humid_Mar Month Temp_F
##   <chr>      <chr>      <chr>      <chr> <chr>
## 1 New York  75%        72%        68%      Jan  32°F
## 2 New York  75%        72%        68%      Feb  35°F
## 3 New York  75%        72%        68%      Mar  42°F
## 4 Los Angeles 65%        63%        60%      Jan  58°F
## 5 Los Angeles 65%        63%        60%      Feb  60°F
## 6 Los Angeles 65%        63%        60%      Mar  65°F

```

```
head(untidy_weather_right)
```

```

## # A tibble: 6 x 6
##   City      Temp_Jan Temp_Feb Temp_Mar Month Humidity_Percent
##   <chr>      <chr>      <chr>      <chr> <chr>
## 1 New York  32°F      35°F      42°F      Jan  75%
## 2 New York  32°F      35°F      42°F      Feb  72%
## 3 New York  32°F      35°F      42°F      Mar  68%
## 4 Los Angeles 58°F      60°F      65°F      Jan  65%
## 5 Los Angeles 58°F      60°F      65°F      Feb  63%
## 6 Los Angeles 58°F      60°F      65°F      Mar  60%

```

```

untidy_weather <- left_join(untidy_weather_left, untidy_weather_right, by = c("City", "Month"))
head(untidy_weather)

```

```

## # A tibble: 6 x 10
##   City      Humid_Jan Humid_Feb Humid_Mar Month Temp_F Temp_Jan Temp_Feb Temp_Mar
##   <chr>      <chr>      <chr>      <chr> <chr> <chr> <chr> <chr> <chr>
## 1 New York  75%        72%        68%      Jan  32°F  32°F    35°F    42°F
## 2 New York  75%        72%        68%      Feb  35°F  32°F    35°F    42°F
## 3 New York  75%        72%        68%      Mar  42°F  32°F    35°F    42°F
## 4 Los Angeles 65%        63%        60%      Jan  58°F  58°F    60°F    65°F
## 5 Los Angeles 65%        63%        60%      Feb  60°F  58°F    60°F    65°F
## 6 Los Angeles 65%        63%        60%      Mar  65°F  58°F    60°F    65°F
## # i 1 more variable: Humidity_Percent <chr>

```

We now have our individual columns for Month, Temperature and Humidity, however we must trim the fat and get rid of each column that has a single months values.

```

untidy_weather <- untidy_weather |>
  select("City",
         "Month",
         "Temp_F",
         "Humidity_Percent")
head(untidy_weather)

```

```
## # A tibble: 6 x 4
##   City      Month Temp_F Humidity_Percent
##   <chr>      <chr> <chr>   <chr>
## 1 New York   Jan   32°F   75%
## 2 New York   Feb   35°F   72%
## 3 New York   Mar   42°F   68%
## 4 Los Angeles Jan   58°F   65%
## 5 Los Angeles Feb   60°F   63%
## 6 Los Angeles Mar   65°F   60%
```


With this, we have our data tidy, but to make sure we can examine with EDA we must change Temperature and Humidity from character values to integers.

```
tidy_weather <- untidy_weather |>
  mutate(Temp_F = parse_number(Temp_F),
         Humidity_Percent = parse_number(Humidity_Percent))
head(tidy_weather)
```


```
## # A tibble: 6 x 4
##   City      Month Temp_F Humidity_Percent
##   <chr>      <chr> <dbl>         <dbl>
## 1 New York   Jan      32           75
## 2 New York   Feb      35           72
## 3 New York   Mar      42           68
## 4 Los Angeles Jan      58           65
## 5 Los Angeles Feb      60           63
## 6 Los Angeles Mar      65           60
```

Now we can begin EDA.


```
tidy_weather |>
  ggplot(aes(x = City, y = Temp_F)) +
  geom_boxplot()
```




```
tidy_weather |>  
  ggplot(aes(x = City, y = Humidity_Percent)) +  
  geom_boxplot()
```

```
tidy_weather |>  
  ggplot(aes(x = reorder(Month, Temp_F), y = Temp_F)) +  
  geom_boxplot()
```




```
tidy_weather |>  
  ggplot(aes(x = reorder(Month, Temp_F), y = Humidity_Percent)) +  
  geom_boxplot()
```



Here we have graphs showing temperature and humidity broken down by either city or month. These graphs show that Chicago was the coldest across sampled months while LA was the hottest, while the opposite is true regarding humidity. These graphs also show that temperature increases as months go on while humidity decreases.

We can also look at how temperature and humidity fluctuate over these months across different cities with the following graph.

```
tidy_weather |>
  ggplot(aes(x = reorder(Month, Temp_F))) +
    geom_point(aes(y = Temp_F, color = "Temp_F"), size = 3) +
    geom_point(aes(y = Humidity_Percent, color = "Humidity_Percent"), size = 3) +
    facet_wrap(~ factor(City, levels = c("Chicago", "New York", "Los Angeles"))) +
    labs(x = "Month", y = "Value") +
    scale_color_manual(values = c("Temp_F" = "blue", "Humidity_Percent" = "red"))
```



Here we can see that across all cities, as the months continue from January to March, temperature increases and humidity decreases.

Emissions

This is the largest dataset of the three, and needs the most work tidying and cleaning before EDA.

```
head(untidy_emissions)
```

##	Area	Item	Element	Unit				
## 1	Afghanistan	Crop Residues	Direct emissions (N2O)	kilotonnes				
## 2	Afghanistan	Crop Residues	Indirect emissions (N2O)	kilotonnes				
## 3	Afghanistan	Crop Residues	Emissions (N2O)	kilotonnes				
## 4	Afghanistan	Crop Residues	Emissions (CO2eq) from N2O (AR5)	kilotonnes				
## 5	Afghanistan	Crop Residues	Emissions (CO2eq) (AR5)	kilotonnes				
## 6	Afghanistan	Rice Cultivation	Emissions (CH4)	kilotonnes				
##	X2000	X2001	X2002	X2003	X2004	X2005	X2006	X2007
## 1	0.520	0.5267	0.8200	0.9988	0.8225	1.1821	1.0277	1.2426
## 2	0.117	0.1185	0.1845	0.2247	0.1851	0.2660	0.2312	0.2796
## 3	0.637	0.6452	1.0045	1.2235	1.0075	1.4481	1.2589	1.5222
## 4	168.807	170.9884	266.1975	324.2195	266.9995	383.7498	333.6093	403.3749
## 5	168.807	170.9884	266.1975	324.2195	266.9995	383.7498	333.6093	403.3749
## 6	18.200	16.9400	18.9000	20.3000	27.3000	22.4000	22.4000	23.8000
##	X2008	X2009	X2010	X2011	X2012	X2013	X2014	X2015
## 1	0.8869	1.3920	1.2742	1.0321	1.3726	1.4018	1.4584	1.2424
## 2	0.1996	0.3132	0.2867	0.2322	0.3088	0.3154	0.3281	0.2795
## 3	1.0865	1.7051	1.5609	1.2643	1.6815	1.7173	1.7865	1.5220

```
## 4 287.9099 451.8647 413.6467 335.0379 445.5958 455.0727 473.4174 403.3181
## 5 287.9099 451.8647 413.6467 335.0379 445.5958 455.0727 473.4174 403.3181
## 6 26.6000 28.0000 29.1200 29.4000 28.7000 28.7000 30.8000 22.9600
##      X2016      X2017      X2018      X2019      X2020
## 1 1.1940 1.0617 0.8988 1.2176 1.3170
## 2 0.2687 0.2389 0.2022 0.2740 0.2963
## 3 1.4627 1.3005 1.1011 1.4916 1.6133
## 4 387.6130 344.6447 291.7838 395.2689 427.5284
## 5 387.6130 344.6447 291.7838 395.2689 427.5284
## 6 16.6600 15.3233 16.4555 17.8542 20.6577
```

While “Item” and “Element” are extended across a “long” format, years are given their own columns, so we can tidy that first. We can also get rid of “Unit” as every “amt” is measured in kilotons and nothing else.

```
untidy_emissions <- untidy_emissions |>
  pivot_longer(
    cols = starts_with("X"),
    names_to = "Year",
    values_to = "Amt_kt"
  )
untidy_emissions <- untidy_emissions |>
  select(!Unit)
head(untidy_emissions)
```

```
## # A tibble: 6 x 5
##   Area      Item      Element      Year  Amt_kt
##   <chr>    <chr>    <chr>    <chr> <dbl>
## 1 Afghanistan Crop Residues Direct emissions (N20) X2000 0.52
## 2 Afghanistan Crop Residues Direct emissions (N20) X2001 0.527
## 3 Afghanistan Crop Residues Direct emissions (N20) X2002 0.82
## 4 Afghanistan Crop Residues Direct emissions (N20) X2003 0.999
## 5 Afghanistan Crop Residues Direct emissions (N20) X2004 0.822
## 6 Afghanistan Crop Residues Direct emissions (N20) X2005 1.18
```

Now that our data is tidy and in a “long” format, we can clean it up for EDA.

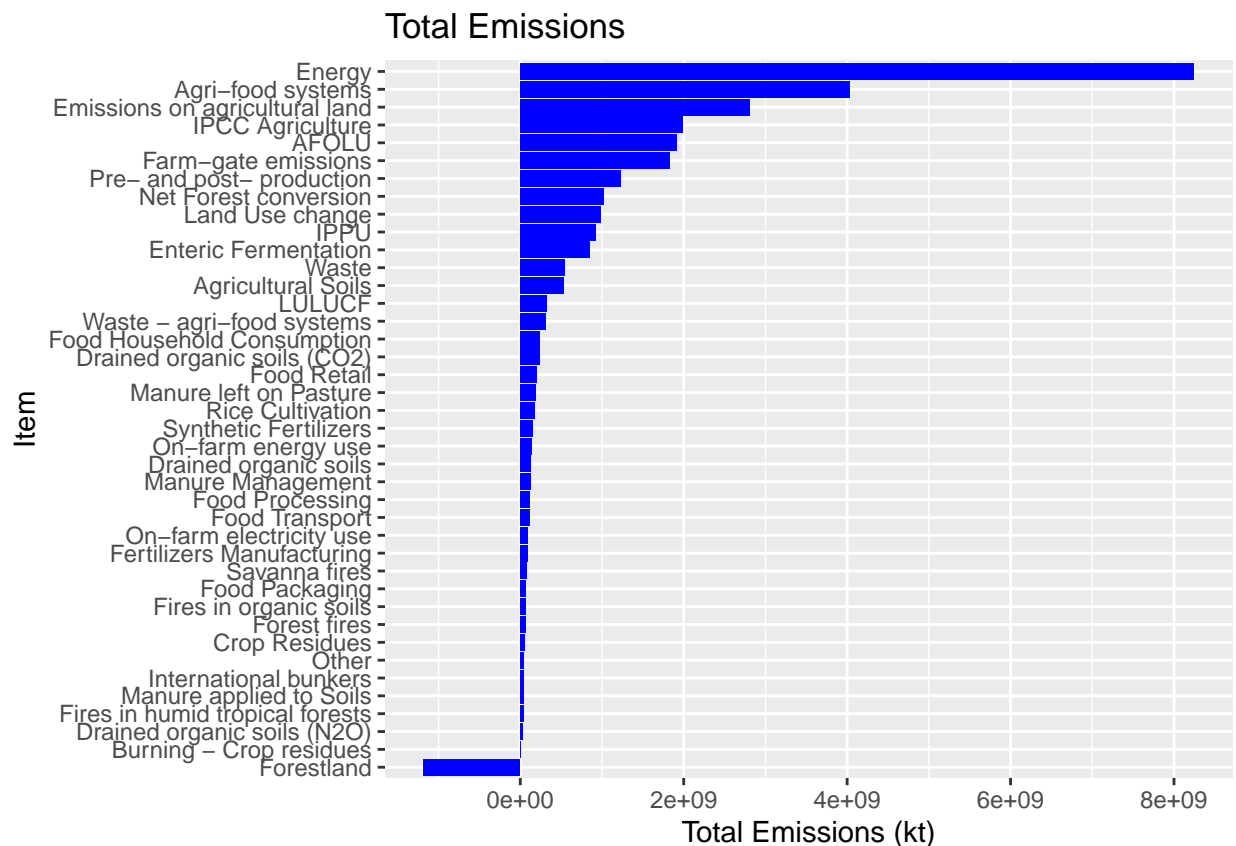
```
tidy_emissions <- untidy_emissions |>
  mutate(Element = str_remove(Element, regex("emissions", ignore_case = TRUE)),
    Element = str_remove(Element, "\\(AR5\\)"),
    Element = str_remove_all(Element, "\\(|\\)"),
    Year = str_remove(Year, "X"))
head(tidy_emissions)
```

```
## # A tibble: 6 x 5
##   Area      Item      Element      Year  Amt_kt
##   <chr>    <chr>    <chr>    <chr> <dbl>
## 1 Afghanistan Crop Residues Direct N20 2000 0.52
## 2 Afghanistan Crop Residues Direct N20 2001 0.527
## 3 Afghanistan Crop Residues Direct N20 2002 0.82
## 4 Afghanistan Crop Residues Direct N20 2003 0.999
## 5 Afghanistan Crop Residues Direct N20 2004 0.822
## 6 Afghanistan Crop Residues Direct N20 2005 1.18
```

With the data tidied and observations cleaned, we can now begin EDA.

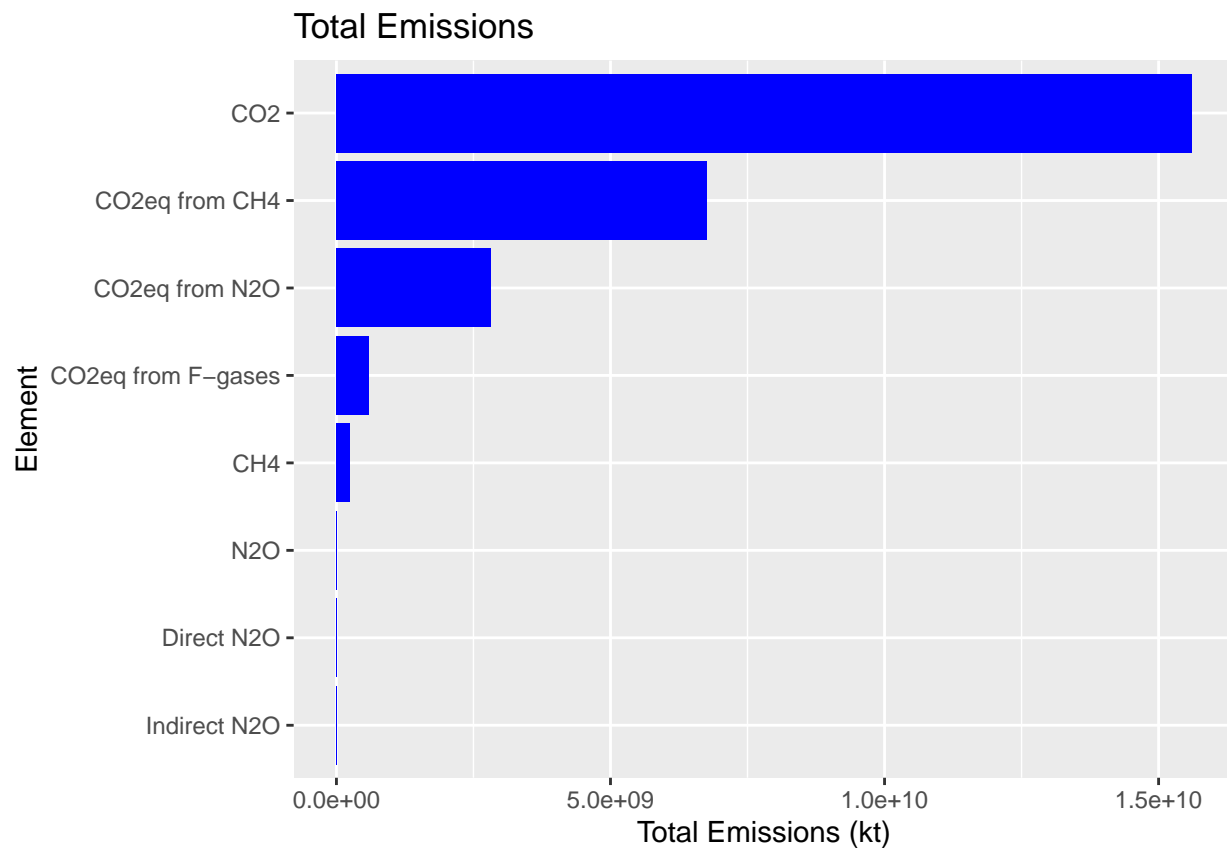
While my classmate suggested looking at overall total emissions per country for each year, and while I do think that would be insightful, processing and presenting that data has proved difficult. Therefore I will look at another suggested potential analysis, that of overall total emissions across source. Additionally, we can look at which element being polluted is the greatest.

```
tidy_emissions |>
  filter(Item != "All sectors with LULUCF" &
         Item != "All sectors without LULUCF") |>
  group_by(Item) |>
  summarise(Total_Emissions = sum(Amt_kt, na.rm = TRUE), .groups = "drop") |>
  ggplot(aes(x = fct_reorder(Item, Total_Emissions, .desc = FALSE), y = Total_Emissions)) +
  geom_col(fill = "blue") +
  labs(title = "Total Emissions",
       x = "Item",
       y = "Total Emissions (kt)") +
  coord_flip()
```




```
tidy_emissions |>
  filter(Element != "CO2eq") |>
  group_by(Element) |>
  summarise(Total_Emissions = sum(Amt_kt, na.rm = TRUE), .groups = "drop") |>
  ggplot(aes(x = fct_reorder(Element, Total_Emissions, .desc = FALSE), y = Total_Emissions)) +
  geom_col(fill = "blue") +
  labs(title = "Total Emissions",
```

```
x = "Element",
y = "Total Emissions (kt)" +
coord_flip()
```



```
tidy_emissions |>
  filter(Element == "N2O" |
         Element == "Direct N2O" |
         Element == "Indirect N2O") |>
  group_by(Element) |>
  summarise(Total_Emissions = sum(Amt_kt, na.rm = TRUE), .groups = "drop") |>
  ggplot(aes(x = fct_reorder(Element, Total_Emissions, .desc = FALSE), y = Total_Emissions)) +
  geom_col(fill = "blue") +
  labs(title = "Total Emissions",
       x = "Element",
       y = "N20 Emissions (kt)") +
  coord_flip()
```



After filtering out variables accounting for all sectors with or without land use change, we find that energy is the greatest driver of emissions. Every item contributes to greater emissions except for one, that being forest land. Looking at the elements graph, we find that the single element emitted the most is CO₂. N₂O polutes the least, however we cannot get a good idea of how much compared to direct and indirect N₂O, so we can break it down even further to find that N₂O is producing a fraction of the other elements. Despite this, its equivalency in CO₂ is the third largest on this list.

Conclusion

While all of these datasets had some stark differences, whether it be by the number of variables, or by the amount of sheer observations contained within, the process for tidying them was roughly the same across all three. Begin by identifying which columns can be addressed as the same variable, extend the data longer with said variable, and then transform the information into something to process. After that is when we see differences, as each dataset could be broken down differently. This leads to different forms of representation being better for one than another - I would not use columns as I did in emissions data to represent temperature and humidity in weather data. This was insightful on how managing and processing data can be very similar and very different!